

Blackjack Reinforcement Learning: An Analysis of Q-Learning Techniques

Ian McAtee and Fletcher Wadsworth

EE5885: Deep Reinforcement Learning and Control

Professor Jiang

University of Wyoming

December 15, 2021

Abstract

This project examines the application of Q-learning techniques to train a reinforcement learning agent to play the game of blackjack. Specifically, the techniques of tabular Q-learning and deep Q-learning are implemented in a simulated blackjack environment. It is shown that while both methods achieve the learning of a blackjack policy, neither technique converged to an optimal policy. Both methods exhibit a blackjack win percentage of approximately 42%, failing to show significant improvement over standard heuristic blackjack policies. A discussion is presented that analyzes the possible reasoning behind these results, the suitability of Q-learning for this application, and examines the difficulties of using reinforcement learning techniques on the game of blackjack.

1 Introduction and Background

Blackjack is one of the most popular casino table games around the world. What differentiates blackjack from other games is the opportunity for strategic choices. Unlike games such as roulette, blackjack is structured so that a player's actions are correlated to their success, such that an optimal strategy can be employed to minimize that player's losses. Certain varieties of casino blackjack have known optimal strategies which are derived statistically and supported empirically [1]. However, applying Reinforcement Learning (RL) techniques to blackjack in an effort to generate an optimal policy experientially is a promising endeavor which aims to understand the underlying dynamics of competitive card games and improve a gambler's chances of winning.

Ultimately, blackjack is a comparatively simple game to which RL has been applied in a variety of implementations. The eventual target application is to the game of poker. Poker is a significantly more complex environment than blackjack; games have a larger state-action space comprised of more cards and available actions and other players are independent agents capable of their own decisions. With the immense number of permutations of a poker hand and the psychology of not entirely rational competitors influencing the trajectory of a hand, current poker strategies are only naively statistically supported and are largely intuitive. Thus, blackjack is the chosen environment used to build a foundation upon which the domain of poker can eventually become the target.

Here, agents are trained towards optimal blackjack strategy in a simulated environment using Q-Learning. Tabular Q-Learning and Deep Q-Learning methods are implemented and their resulting policies are compared in an effort to understand the advantages of each.

2 Methods and Results

2.1 Environment

To reduce the state and action space of the blackjack environment to a manageable level, the following rule-set is used in this implementation:

Blackjack Environment

- ◆ Dealer and player each receive two cards from standard 52 card deck
 - ▷ One of the dealer's cards is face up and known to the player
- ◆ Face cards (Jack, Queen, King) all valued at 10
- ◆ Aces are valued at 11 or 1
- ◆ Player has two available actions at each turn: **hit** and **stay**
- ◆ At each turn, player selects one of the actions and is dealt another card or the hand terminates
- ◆ After the player exceeds 21 (bust) or chooses to stay, the dealer resolves their hand
 - ▷ Dealer follows a fixed policy in which they must hit below 17
- ◆ The winner is determined
 - ▷ If the player's hand exceeds the dealer's hand and the player has not busted, the player wins
 - ▷ If the player busts or the dealer's hand exceeds the player's without busting, the player loses
 - ▷ If neither the player or dealer bust and they have equal hand values, the hand is a draw (push)
- ◆ A terminal state is reached and a reward is administered to the player based on the outcome of the hand

From this environment, the state, action, and reward spaces can be formalized as

$$S = \{sum, up\ card, usable\ ace\} \tag{1}$$

$$A = \{hit, stay\} \tag{2}$$

$$R = \{-wager, 0, wager\}, \tag{3}$$

where *sum* is the sum of the cards in the player's hand, *up card* is the visible dealer card, *usable ace* is whether the player hand contains an ace which can be reduced to 1, and *wager* is the bet amount won or lost. Note that the player agent has no control over the wager amount, making the reward a design parameter rather than a learned action.

2.2 Tabular Q-Learning

In Q-Learning, $Q(s, a)$ is known as the *action-value* function. $Q(s, a)$ outputs the value of a particular action in a particular state. The action-value function is formalized as

$$Q^\pi(s, a) = E[G_t | s_t = s, a_t = a] \quad (4)$$

where

$$G_t = r_t + \gamma r_{t+1} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k}. \quad (5)$$

Thus, $Q(s, a)$ is the expected reward of executing an action a in the state s . In the tabular Q-Learning process, $Q(s, a)$ is not an analytical function, but rather a table which is populated and updated during learning and has a well defined output for each state-action pair. The values in this Q-table are learned using the update rule

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_a Q(s', a) - Q(s, a)] \quad (6)$$

where α is the learning rate, γ is the discount factor, and s' is the next state.

This update rule implies Q-Learning as an off-policy method; updates are made using a *greedy* policy by selecting the most rewarding action in the given next state, while the action a taken to reach the next state need not be chosen via a greedy policy. The policy used to select actions during exploration is known as the *behavioral* policy.

Q-learning can also be implemented offline. Offline RL denotes a learning process by which the static data set (which consist of $\{s, a, r, s'\}$ tuples) can be obtained before training and sampled during training. Such an offline implementation is utilized here with 3,000,000 simulated hands, and the training process is summarized:

Tabular Q-Learning

1. Dataset $D = \{s_i, a_i, r_i, s'_i\}$, $i = 1, 2, \dots, N$ is obtained by rollout of a behavioral agent
 - ▷ Agent follows a random uniform policy, i.e. hit or stay with probability 0.5
2. Q-table initialized to 0
3. A single data tuple is sampled from D
4. Q-table entry is updated by Equation 6
5. Repeat 3,4 for all tuples in D

After the above learning procedure is complete, the learned policy is instantiated by a greedy query of the Q-table

$$\pi(s) = \max_a Q(s, a) \quad (7)$$

Once agent is trained on the data set, an evaluation of the agent's performance is performed over 200,000 hands. The win percentage of the trained agent is compared with that of one agent following a uniform random policy and one following a heuristic policy. The heuristic policy is to hit when the hand is valued below 17 and stay otherwise.

Examples of learned greedy policies are shown below, where the action with maximum value for each state is indicated. In these policy tables, red and blue denote *hit* and *stay* as the preferred action respectively:

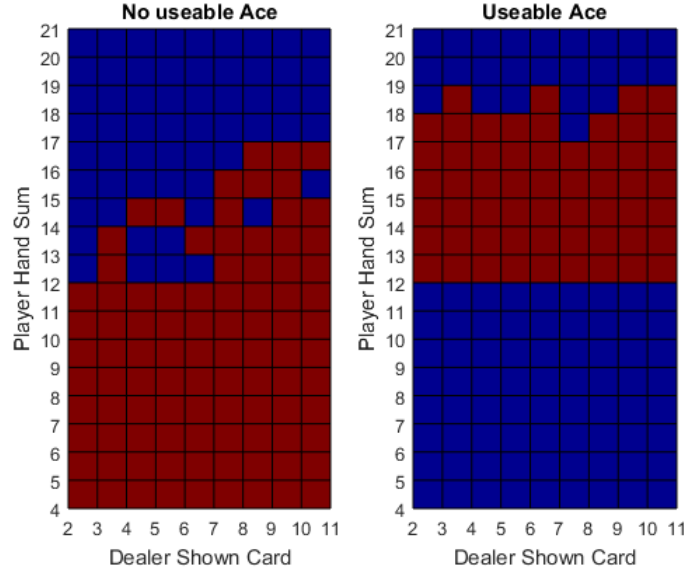


Figure 1: Greedy Policy with $\alpha = 0.01$ and $\gamma = 0.95$

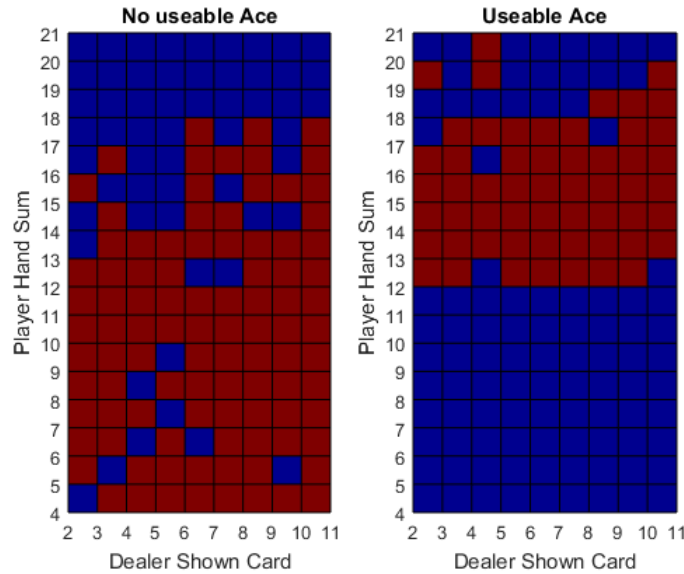


Figure 2: Greedy Policy with $\alpha = 0.3$ and $\gamma = 0.6$ (Red and Blue Denote Hit and Stay, Respectively)

The following tables show the performance evaluation response for the corresponding learned policies:

| | |
|---------------------------|--------|
| Trained Agent Win Rate | 0.4184 |
| Random Policy Win Rate | 0.3042 |
| Heuristic Policy Win Rate | 0.4243 |

Table 1: Greedy Policy Performance with $\alpha = 0.01$ and $\gamma = 0.95$

| | |
|---------------------------|--------|
| Trained Agent Win Rate | 0.4172 |
| Random Policy Win Rate | 0.3042 |
| Heuristic Policy Win Rate | 0.4223 |

Table 2: Greedy Policy Performance with $\alpha = 0.3$ and $\gamma = 0.6$

2.3 Deep Q-Learning

Having demonstrated the application of tabular Q-learning to the game of blackjack, deep Q-learning was posited as a possible method to improve results. Deep Q-learning is a neural network parameterized extension of the classical Q-learning algorithm. Recall that in tabular Q-learning, one must catalogue the Q-function, meaning that each action-state value pair must be tabulated, tracked, and updated throughout the course of learning. This is manageable given a limited and discrete action-state space, however, becomes increasingly intractable as this space expands. Deep Q-learning seeks to solve this intractability by using neural networks to approximate the Q-function of a state, rather than attempt to store and index an immense Q-table. The crux of deep Q-learning is the idea that a neural network can be trained as to accurately parameterize the Q-function such that one can obtain an estimate of the action-value for all actions given the state. This concept is shown below in figure 3.




Figure 3: Overview of the Deep Q-Network Implementation

It should be noted that strictly speaking, deep reinforcement learning approaches are not essential for the solution of the blackjack problem. Recall that here, the state and action space, while large, have been shown to be tractable via tabular Q-learning methods. However, the goal of the application of deep Q-learning to the blackjack problem is to comprehend its implementation, compare the results against the tabular method, and to discuss its suitability as a possible solution to more complex card games such as poker.

2.3.1 Blackjack Environment Modifications

To apply a deep Q-learning approach to the blackjack problem, modifications had to be made to the environment as created for the tabular Q-learning method. In order to facilitate the straightforward application of the neural network parameterization, a prebuilt blackjack environment was imported from the *OpenAI Gym* library. This environment is of the same form as presented in section 2.1. However, the state space and reward function were modified slightly as shown:

$$S = \{sum, up\ card\} \quad (8)$$

$$A = \{hit, stay\} \quad (9)$$

$$R = \{-1, 0, 1.5\} \quad (10)$$

Here, one can see that the state space S was truncated as to not include the *usable ace* state. A result of the elimination of the *usable ace* is that the value of the ace was set at one for this instance of the blackjack environment. The state space was changed as to simplify its presentation to the associated Q-learning networks. Because the goal of the deep Q-learning implementation was to comprehend its operation, the networks were constructed from the ground-up. Thus, the state space reduction was performed purely to lessen the complexities of the neural network design process. It is recommended that future work in the application of deep Q-learning to blackjack include the *usable ace* state for the purposes of accuracy and completeness.

It can also be noticed that in equation 10, the reward function was also adjusted. This modification was made in attempt to facilitate better learning of the agent in the blackjack environment. The underlying belief is that because blackjack is inherently biased against the player, that is, the player is statistically predisposed to losing, the reward should be higher for a win as compared to a loss. However, this reward function still suffers from the issue of sparsity in time due to zero reward being yielded until the termination of the blackjack hand. However, a more robust discussion of the reward design process and its intrinsic difficulties in blackjack is given in section 3.2.

2.3.2 Deep Q-Learning Methods

The blackjack environment used allowed for the implementation of an *online* method of deep Q-learning. That is, the environment yields a training tuple $\{S, A, S', R\}$ at every time step. The pseudocode of the *online* deep Q-learning implementation is given below and a block-diagram overview is shown in figure 4.

Online Deep Q-Learning

1. Take ϵ -greedy action a_i to obtain the tuple sample $\{s_i, a_i, s'_i, r_i\}$ from environment
2. Compute TD target as: $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'} Q(s'_i, a'_i; w')$
 - $Q(s'_i, a'_i; w')$ is computed using **target network**
3. Update the **Q-function approximation network** weights
 - $w \leftarrow w - \alpha \sum_i \Delta_w Q(s_i, a_i; w)(Q(s_i, a_i; w) - y_i)$
 - $Q(s_i, a_i; w)$ is computed using **Q-function approximation network**
4. Update **target network** w' with **Q-function approximation network** w every M steps
5. Repeat 1-4 for all training tuples

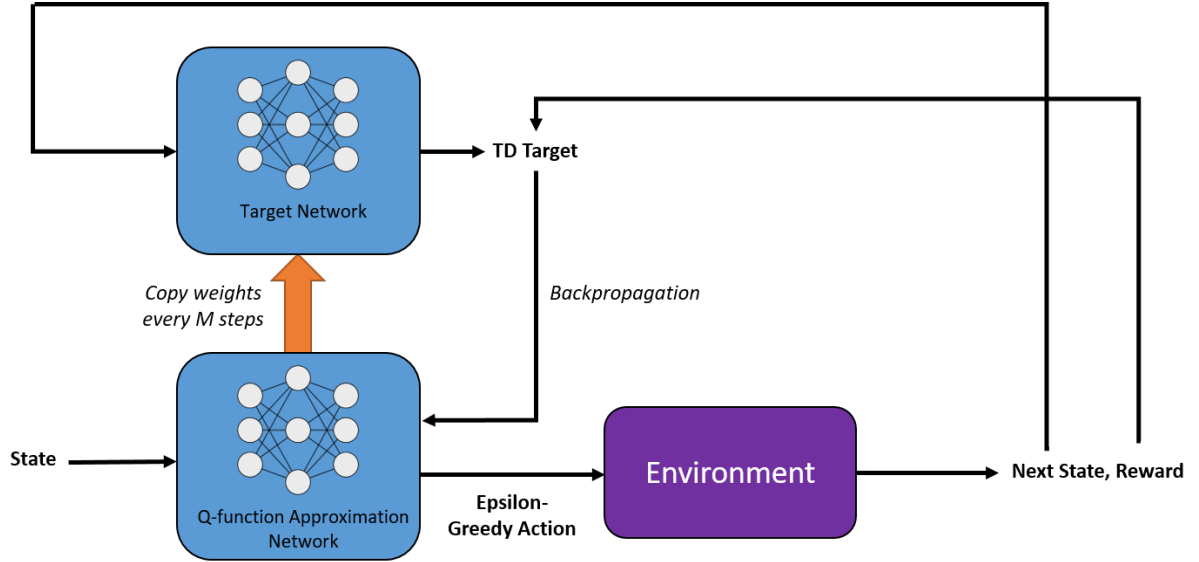


Figure 4: Overview of Deep Q-learning Network Training

As one can see, this implementation involves the creation of dual neural networks. The *Q-function approximation network* is the neural network parameterization of the Q-function. It is input with the state space and provides an estimation of the action value of each of the possible actions in the action space. Here, an ϵ -greedy policy is used to select the action taken by the agent in the environment. That is, the agent will with a $1 - \epsilon$ probability, take the greediest action. This greedy action is the one associated with the maximal action-value estimation. It is this Q-function approximation network that one wishes to optimize through learning such that it is an accurate representation of the Q-function and thus can yield a satisfactory policy.

However, the training of this network can prove difficult due to challenges inherent in Q-learning. These challenges are well-known and studied and include issues of nonstationarity and sample correlation. If one examines equation 6, it can be noticed that the *temporal difference* (TD) target $r + \max_a Q(s', a)$ is not stationary. That is, the target continuously changes upon each iteration. Recall that in normal deep learning neural networks, the target is stationary and thus the training is stable. Consequently, this non-stationary target in deep Q-learning means that training can exhibit highly divergent behavior. The solution to this is to introduce a separate but identically architected network to calculate the target value. This target network exhibits set parameters that are only updated every M training steps as a cloning of the Q-function approximation network parameters. The rational is that by fixing the target values, the training becomes pseudo-stationary within the M -step training window.

The issue of sample correlation arises due to the online nature of the deep Q-learning algorithm. The sample tuples are generated from the environment sequentially, meaning that in general, they are highly correlated. To account for this, the Q-learning algorithm usually implements experience replay. In experience replay, generated samples are stored in a replay buffer and batched samples are randomly drawn from the buffer to perform the learning. This random drawing of the samples provides for the decorrelation of the training samples. However, the highly stochastic nature of blackjack means that sequence of state transitions does not happen in a linear and organized way. Therefore, the assumption is made that the samples generated by the blackjack environment are not particularly correlated and thus, the implementation of experience replay is not necessary. A deeper discussion of this assumption is presented in section 3.1.

2.3.3 Deep Q-Learning Results

The deep Q-learning method described in section 2.3.2 was applied to the modified blackjack environment. Through experimentation, a simple two hidden layer fully connected neural network was found to be sufficient for blackjack Q-function approximation. The specifics of the network architecture used for both the Q-function approximation network and the target network is given in table 3.

| Layer | Number of Nodes | Activation Function |
|----------|-----------------|---------------------|
| Input | 2 | Identity |
| Hidden 1 | 32 | RELU |
| Hidden 2 | 16 | RELU |
| Output | 2 | SoftMax |

Table 3: Neural Network Architecture of Q-Function Approximation and Target Networks

Training of the Q-function approximation network was accomplished by implementing deep Q-learning with 2500 samples generated by the blackjack environment in an online paradigm. The following parameters in table 4 were used to implement the deep Q-learning:

| Parameter | Value |
|------------------------------------|-----------|
| Exploration Rate ϵ | 0.1 |
| Discount Factor γ | 0.9 |
| Learning Rate α | 0.001 |
| Target Network Update Interval M | 100 Steps |

Table 4: Deep Q-Learning Parameters

The 2500 samples were organized into 50 episodes, each with 50 samples, in an attempt to analyze how the average reward gained by the agent changed over time. Figure 5 plots this average reward versus the number of episodes.




Figure 5: Average Reward Over Training Episodes

An examination of figure 5 shows that the deep Q-learning was unable to produce an upward trend in the average rewards. Instead, it is observed that the average reward steadily fluctuates around zero, with the average reward of most episodes remaining negative. This demonstrates the difficulties inherent to blackjack. The fact that the game is biased against the player will make it difficult for any reinforcement learning agent to obtain a steady increase in average rewards and terminate with a policy that will consistently yield positive rewards.

To analyze the resultant blackjack policy achieved by deep Q-learning, the trained Q-function approximation network was used to evaluate the Q-function over the entire state space. The maximum action-value was then used to determine the action the policy would implement in each particular state. This resulted in the policy shown in figure 6.




Figure 6: Blackjack Policy Learned by Deep Q-Learning (Red and Blue Denote Hit and Stay, Respectively)

Figure 6 shows that the policy found by deep Q-learning is much more uniform than that achieved by tabular Q-learning. This is likely the result of the neural network approximation of the Q-function. Unlike tabular methods, the deep approach means that the neural network can generalize, even if a particular state hasn't been experienced often. Therefore, one would expect a more uniform policy. However, this does not necessarily mean the policy is superior to the tabular policy. This generalization may mean that subtleties in the state-transition statistics have been lost that may be relevant in defining a truly optimal policy. In fact, the policy produced by deep learning is strikingly similar to that of the heuristic *hit under 17* policy, which is known to not be the optimal policy. Therefore, it is unsurprising that upon testing, the deep Q-learning policy and the heuristic policy exhibit similar win percentages, as shown in table 5.

| Policy | Win Percentage |
|-----------------|----------------|
| Heuristic | 0.428 |
| Deep Q-Learning | 0.425 |

Table 5: Win Percentages of Heuristic and Deep Q-Learning Policies as Evaluated Over 1000 Blackjack Rounds

3 Discussion

3.1 State Transition Stochasticity

The win percentage of the trained agents demonstrate that the agents have learned a policy sufficient to outperform a random policy. While this is promising, the inability of the agents to significantly outperform the heuristic policy is worthy of exploration.

Consider the following region of the state space in which the player has no usable aces:




Figure 7: State Space Region of Interest

This region represents a departure from the heuristic policy which can be reasonably analyzed. Since all face cards are valued at 10, the plurality of the deck is valued at 10. Assuming multiple decks are combined (commonplace in casino blackjack), regardless of the cards dealt to the dealer and players, there is an approximate 30.8% chance that the dealer has a 10 card hidden, and the same chance of the player or the dealer drawing a 10 card. This means that in this region, both the player and dealer are more likely to bust than not. However, the dealer is beholden to the policy placed on them, and must hit for hands less than 17. From this it can be surmised that despite a sub-optimal player hand (not valued high enough for a reasonable chance of beating another hand), the dealer is likely to bust, and thus the player should stay.

However, this is a rather subtle advantage for the player. Since the player does not know the entire hand of the dealer, and cards are drawn randomly, the outcome of this particular premise has high variance. To testify to this point, consider the following state transition:

| | | | | | | | | | | | |
|-----------------|----------------|---|---|---|---|-----|---|---|---|----|----|
| Player Hand Sum | Bust | | | | | | | | | | |
| | 21 | | | | | | | | | | |
| | 20 | | | | | | | | | | |
| | 19 | | | | | 8% | | | | | |
| | 18 | | | | | 31% | | | | | |
| | 17 | | | | | 8% | | | | | |
| | 16 | | | | | 8% | | | | | |
| | 15 | | | | | 8% | | | | | |
| | 14 | | | | | 8% | | | | | |
| | 13 | | | | | 8% | | | | | |
| | 12 | | | | | 8% | | | | | |
| | 11 | | | | | 8% | | | | | |
| | 10 | | | | | 8% | | | | | |
| | 9 | | | | | | | | | | |
| | 8 | | | | | | | | | | |
| | 7 | | | | | | | | | | |
| | 6 | | | | | | | | | | |
| | 5 | | | | | | | | | | |
| | 4 | | | | | | | | | | |
| | | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |
| | Dealer Up Card | | | | | | | | | | |

| | |
|--|----------------------|
| | Current State |
| | Possible Next States |

Figure 8: State Transition Probability

In any given state, the number of states to which can be transitioned by choosing to hit is at most 10. Furthermore, the distribution of the possible next state has fairly high variance. These examples illustrate the profound stochasticity of the state transitions inherent in card games such as blackjack. As a consequence, an enormous amount of experiential data is needed for the agent to learn the marginal advantage that one action may possess over another in a given state which does not appear naturally or obviously.

3.2 Reward Design

The reward schedule which comes intuitively from the structure of blackjack is to tender a reward when the hand has been resolved and a winner decided and tender no reward for agent actions during the episode. This is a natural reward scheme, since in an actual game of blackjack, a player wins no money and is not compensated or penalized for taking actions per se. This leads to some issues with regard to the Q-table and network parameter updates with regard to tabular and deep Q-Learning respectively. Consider Equation 6; with the current reward schedule, every $\{s, a, r, s'\}$ tuple collected has the form of $\{s, a, 0, s'\}$ or $\{s, a, r, s_{terminal}\}$. Thus every update contains a reward term or a discounted next state greedy action-value, but not both. Such an environment can be characterized as a sparsely rewarded environment. Sparse rewards make updating Q-table values difficult, particularly at the beginning of learning, as the discounted next state action-value may not be populated or contain meaningful information. Similarly for the parameter update in deep Q-Learning, initial learning with sparse rewards results in poor initial learning and is subject to the forces

of random weight initialization and uncorrelated data.

A simple and obvious strategy to overcome these sparse reward difficulties is more training time and a larger training data set. This is a through line to card game problems in RL as has been discussed; the narrow statistical advantages throughout the state space and sparse reward schedule benefit largely from expansive training. However, a redesign of the reward schema could be a viable tactic to combat the need for ever-larger training time. While the implementations here utilized the intuitive reward scheme discussed, alternatives were considered. In particular, different scaling of the wager amount showed no notable effect on the learned policy. Other potential renovations to the reward schedule include different ratios of reward for different outcomes (i.e. change the ratio of punishing a losing hand to rewarding a winning one), a dynamic reward schedule where the reward amount would either decay or increase as training progresses, and adding rewards to actions taken regardless of the outcome of the episode. It can also be noted that a different RL technique entirely could be effectively utilized to attenuate this problem. Monte-Carlo Learning is a possible alternative to Q-Learning, as it learns directly from sampled trajectories. Revising of the reward schema and exploration of other RL techniques are future endeavors to further develop this problem.

4 Conclusion and Future Work

This paper analyzed blackjack as an application for reinforcement learning. In particular, tabular and deep Q-Learning were used to train a blackjack agent in hopes of reaching an optimal policy, wherein the trained agent could minimize its losses in a series of simulated games and perform better than an agent following a more heuristic but less robust policy. In these implementations, the trained agents performed about on par with the heuristic policy and failed to achieve an optimal policy. Several observations have been made about possible reasons for the underwhelming trained policy, including the profoundly stochastic nature of card games, the need for a considerable amount of training time and data, and the difficulty of designing a proper reward schedule.

In the future, fine tuning of the learning hyperparameters and network structure may help to improve the agent's policies, particularly in the deep Q-Learning domain. Exploration of different reward frameworks and perhaps different RL algorithms all together may be more suited to deriving an optimal policy in the context of blackjack and other card games. After these problems are addressed, the expansion of the state and action space to equate the environment to proper casino blackjack and eventual tailoring to more complex card games such as poker is endeavored.

5 Appendix

5.1 Tabular Q-Learning Listings

- Generate Data Script

```
1 %Script generates {s,a,r,s'} data tuples for Q-learning of
  %2 player blackjack environment using uniform random
    [+]behavioral policy
3 clear all
  tic
5 %ACTION: [1,2] == ['stay','hit']
  numEpisodes = 3000;
7 handsPerEpisode = 1000;

9 ante = 10; %reward amount
  numPlayers = 2; %Agent and dealer
11 numCards = 2; %Blackjack initial hand = 2 cards
  numDecks = 4;
13 aceHigh = 1;

15 encode_state = @(hand,dCard,ace) (hand-1) + 18*(dCard-2)+
    [+]180*(ace);

17 SARS = int16([]);
  SARS_total = int16([]);
19
  for i = 1:numEpisodes
21     rng('shuffle')
      SARS=int16([]);
23     parfor ii = 1:handsPerEpisode
        handOver = 0;
25         while ~handOver
            %-----Params and storage arrays-----
27             state = [];
              state_fcn = [];
29             action = [];
              reward = [];

31             %-----Simulate Hand-----
33             %Deal cards to numCards x numPlayers matrix
              [dealtCards,deck] = dealCards(numPlayers,numCards,
                [+]numDecks,aceHigh);
35             %Split cards into player and dealer hands
              pHand = dealtCards(1,:); dHand = dealtCards(2,:);
37             %Evaluate player hand
```

```

39         [pHand, pVal, pBust, pUseAce] = HandEval(pHand);

41     %Select action using random policy
    while (~pBust)&&(~handOver)
43         state = [state, [pVal; dHand(2); pUseAce]];
        state_fcn = [state_fcn, encode_state(pVal,
45             [+]dHand(2), pUseAce)];
        reward = [reward 0];
        actSelect = rand;
        if actSelect < 0.5
47             action = [action, 2];
            [pHand, deck] = PlayerHit(pHand, deck);
49             [pHand, pVal, pBust, pUseAce] = HandEval(
                [+]pHand);
        else
51             action = [action, 1];
            handOver = 1;
53             break
        end
55     end
    %Resolve dealer hand
57     [dBust, dBJ, dVal, deck] = DealerPolicy(dHand, deck);
    %Determine winner
59     if (pBust) || ((dVal > pVal) && (~dBust))
        reward = [reward, -ante];
61         %Lose state encoded by 0
        state_fcn = [state_fcn, 0];
63     elseif (pVal > dVal) || dBust
        reward = [reward, ante/5];
65         %Win state encoded by 2
        state_fcn = [state_fcn, 2];
67     else
        reward = [reward, 0];
69         %Draw state encoded by 1
        state_fcn = [state_fcn, 1];
71     end
        reward(1) = [];
73     for p = 1:length(action)
        SARS = [SARS; [state_fcn(p), action(p),
            [+]reward(p), state_fcn(p+1)]];
75     end
    end
77     SARS_total = [SARS_total; SARS];
79 end

```



```

toc
81 save( 'SARS_data.mat' , 'SARS_total' , '-mat' )

• Deal Cards Function

1 function [dealtCards,deck] = dealCards(numHands, cardsPerHand,
    [+] numDecks, aceHigh)
    %function deals cards from specified number of decks for input
    [+] number of
3 %cards to input number of players
    deck = [];
5    one = ones(1,4); %1x4 vector of ones
    cards = zeros(cardsPerHand,numHands);
7    if aceHigh
        c_set = [2 3 4 5 6 7 8 9 10 10 10 10 11];
9    else
        c_set = [1 2 3 4 5 6 7 8 9 10 10 10 10];
11    end
    %Repeat for each 52 card deck desired in complete shoe
13    for d = 1:numDecks
        %Repeat for each number 1–10 plus 3 face cards (Ace
        [+] encoded as 1)
15        for c = 1:4
            %Append four cards of given number to deck, one for
            [+] each suit
17            deck = [deck; c_set];
        end
19    end
    %make deck into a row vector
21    deck_vec = reshape(deck,1,[]);
    %Shuffle deck
23    deck_vec = deck_vec(randperm(length(deck_vec)));
    %Repeat for desired number of cards per hand
25    for cc = 1:cardsPerHand
        %Repeat for number of players
27        for p = 1:numHands
            %Deal final card in deck
29            cards(p,cc) = deck_vec(end);
            %remove dealt card from deck
31            deck_vec(end) = [];
        end
33    end
    dealtCards = cards;
35    deck = deck_vec;

```

37 **end**

- Player Hand Evaluation Function

```
1 function [newHand, value, bust, useableAce] = HandEval(hand)
   %function evaluates current player hand
3   %
   %Outputs:
5   %
   % newHand: hand with face cards and aces resolved
7   % value: sum of current hand
   % bust: flag indicating if hand value exceeds 21
9   % useableAce: flag indicating whether hand has ace of value
   [+]11

11   %Determine aces in hand
   heldAces = (hand==11);
13   numAces = sum(heldAces);

15   %If no aces in hand
   if numAces == 0
17       %Find hand value
       value = sum(hand);
19       useableAce = 0;
       if value <= 21
21           bust = 0;
       else
23           bust = 1;
       end

25   %If at least one Ace in hand
27   else
       %Find hand value
       value = sum(hand);
29       %if not bust
       if value <= 21
31           useableAce = 1;
           bust = 0;
33       else
           aceCount = numAces;
           value_temp = value;
37           while value_temp > 21 && aceCount > 0
               value_temp = value_temp - 10;
               aceCount = aceCount - 1;
39           end
       end
   end
```

```

41         if value_temp > 21
42             useableAce = 0;
43             bust = 1;
44         else
45             a = find(hand==11);
46             k = 1;
47             while value > 21
48                 hand(a(k)) = 1;
49                 value = sum(hand);
50                 k = k+1;
51             end
52             bust = 0;
53             if isempty(find(hand==11))
54                 useableAce = 0;
55             else
56                 useableAce = 1;
57             end
58         end
59     end
60 end
61 newHand = hand;
end

```

- Player Hit Function

```

function [newHand,newDeck] = PlayerHit(hand,deck)
2 %Function administers card to input hand from input deck
    card = deck(end);
4    deck(end) = [];
    newHand = [hand, card];
6    newDeck = deck;
end

```

- Dealer Policy Function

```

function [bust,blackjack,handValue,newDeck] = DealerPolicy(
    [+]hand, deck)
2
4
    %Obtain deck size (13k x 4)
6    deckSize = size(deck);
8
    %Initialize flags and counter
    blackjack = 0;
10    bust = 0; % Control flag to indicate hand score above 21

```

```

12 hit = 1; % Control flag to indicate hand score below 17
   cardCount = 2; % Counter of hand size
   Hand = hand; % Initial dealer hand vector without suits
14
   %Repeat process as long as dealer has not busted and hand
   [+]value is
16 %below 17
   while ~bust && hit
18       %Hand(Hand > 10) = 10; % Replace face cards with value
       [+] of 10
       heldAces = (Hand == 1); % Indices of aces in dealer
       [+]hand
20       aces = sum(heldAces); % Number of aces in dealer hand
       %Evaluate for no aces in hand
22       if aces == 0
           handValue = sum(Hand); % Value of hand
24
           if handValue > 21 % Dealer bust
26               bust = 1;
           elseif handValue > 16 % Dealer must stay above 17
28               bust = 0;
               hit = 0;
30           else
               card = deck(end);
32               deck(end) = [];
               %Append drawn card to hand vector
34               Hand = [Hand, card];
           end
36       %At least one Ace in dealer hand
       else
38           Hand(heldAces) = 11;
           handValue = sum(Hand);
40           %Reduce soft hands, use ace as 1 instead of 11
           while handValue > 21 && aces > 0
42               handValue = handValue - 10;
               aces = aces - 1;
44           end
           if handValue > 21 %Dealer bust
46               bust = 1;
           elseif handValue > 16 %Dealer must stay above 17
48               bust = 0;
               hit = 0;
50           else
               card = deck(end);
52               deck(end) = [];

```

```

54         %Append drawn card to hand vector
        Hand = [Hand, card];
56     end
58 end
59 newDeck = deck;
    handValue = sum(Hand);
60 end

```

- State Feature Decoding Function

```

1 function [ state ] = StateFunction2Table( enc_state )
    %Function decodes encoded state value (scalar) into individual
3 %state features (1x3 vector)
    %
5 %   encoded_state = (sum - 1) + 18(dealerCard -2) + 180(
    [+])useableAce)
    %
7 %   Output: state = [sum,dealerCard,useableAce]
    orig_enc_state = enc_state;
9     if enc_state < 183
        ace = 0;
11    else
        ace = 1;
13        enc_state = enc_state - 180;
    end
15
    d = mod(enc_state,18)+1;
17
    if d >= 4
19        sum = d;
    else
21        sum = 18+d;
    end
23    dealerCard = (orig_enc_state - 180*ace - (sum-1))/18 + 2;
    state = [sum,dealerCard,ace];
25 end

```

- Q-Learning Script

```

%Q-learning script
2 clear all
    %Load data
4 import = load('SARS_data.mat');
    SARS = import.SARS_total;
6 SARS = single(SARS);

```

```

%State feature vectors
8 sums = 4:21;
  shownCards = 2:11;
10 aces = 1:2;
  actions = 1:2;
12 %Hyperparameters
  alpha = 0.3;
14 gamma = 0.6;
  %Check imported data array
16 [data_samples, sample_size] = size(SARS)
  if sample_size ~= 4
18     error("Data points in SARS not correct dimensions")
  end
20
  %———Parse Data into state, action, reward, and next state
  [%]vectors———
22 %Decode states from integer encoding to state indices for Q
  [%]table
  s = zeros(data_samples,3);
24 sp = zeros(data_samples,3);
  for n = 1:data_samples
26     s(n,:) = StateFunction2Table(SARS(n,1));
     sp(n,:) = StateFunction2Table(SARS(n,4));
28 end
  %Slice action and reward vectors
30 a = SARS(:,2); r = SARS(:,3);
  %Change encoding of useable ace state component for
32 %use as index
  s(:,3) = s(:,3)+1;
34 sp(:,3) = sp(:,3)+1;
  %Initialize Q-table and state visit counter
36 Q_table = zeros(sums(end),shownCards(end),aces(end),actions(
  [%]end));
  state_visit = zeros(sums(end),shownCards(end),aces(end));
38 %Update Q-table
  for k = 1:data_samples
40     maxQ = max(Q_table(sp(k,1),sp(k,2),sp(k,3),:));
     Q_table(s(k,1),s(k,2),s(k,3),a(k)) = Q_table(s(k,1),s(k,2)
     [%],s(k,3),a(k)) ...
42     + alpha*(r(k) + gamma*maxQ - Q_table(s(k,1),s(k,2),s(k
     [%],3),a(k)));
     state_visit(s(k,1),s(k,2),s(k,3)) = state_visit(s(k,1),s(k
     [%],2),s(k,3))+1;
44 end

```

```

46 %Save trained Q-table
save('TrainedQTable.mat','Q_table','-mat')
48
%-----Plotting-----
50 Q_plot = zeros(sums(end),shownCards(end),aces(end));

52 for ii = 4:sums(end)
    for jj = 2:shownCards(end)
54         for kk = aces
            if Q_table(ii,jj,kk,2) > Q_table(ii,jj,kk,1)
56                 Q_plot(ii,jj,kk) = 2;
            else
58                 Q_plot(ii,jj,kk) = -2;
            end
60         end
    end
62 end

64 figure(1)
subplot(1,2,1)
66 surf(Q_plot(:, :, 1))
axis([2 11 4 21])
68 xticks(2:11)
yticks(4:21)
70 colormap('jet')
title('No useable Ace')
72 view(0,90)
ylabel('Player Hand Sum')
74 xlabel('Dealer Shown Card')

76
subplot(1,2,2)
78 surf(Q_plot(:, :, 2))
axis([2 11 4 21])
80 xticks(2:11)
yticks(4:21)
82 colormap('jet')
title('Useable Ace')
84 view(0,90)
ylabel('Player Hand Sum')
86 xlabel('Dealer Shown Card')

88 figure(2)
subplot(1,2,1)
90 surf(Q_table(:, :, 1, 1))

```

```

axis([2 11 4 21])
92 xticks(2:11)
   yticks(4:21)
94 colormap('jet')
   title('Stay Value')
96 view(0,90)
   ylabel('Player Hand Sum')
98 xlabel('Dealer Shown Card')

100
subplot(1,2,2)
102 surf(Q_table(:, :, 1, 2))
   axis([2 11 4 21])
104 xticks(2:11)
   yticks(4:21)
106 colormap('jet')
   title('Hit Value')
108 view(0,90)
   ylabel('Player Hand Sum')
110 xlabel('Dealer Shown Card')

112 figure(3)
   subplot(1,2,1)
114 surf(state_visit(:, :, 1))
   axis([2 11 4 21])
116 xticks(2:11)
   yticks(4:21)
118 colormap('jet')
   title('Ace')
120 view(0,90)
   ylabel('Player Hand Sum')
122 xlabel('Dealer Shown Card')

124
subplot(1,2,2)
126 surf(state_visit(:, :, 2))
   axis([2 11 4 21])
128 xticks(2:11)
   yticks(4:21)
130 colormap('jet')
   title('No Ace')
132 view(0,90)
   ylabel('Player Hand Sum')
134 xlabel('Dealer Shown Card')

```


- Trained Agent Comparison Simulation Script

```

%Script to run games to test agent performance
2  rng( 'shuffle' )
   import = load( 'TrainedQTable.mat', '-mat' );
4  Q = import.Q_table;

6  numPlayers = 2; %Agent and dealer
   numCards = 2; %Blackjack initial hand = 2 cards
8  numDecks = 4; %number of decks in shoe
   aceHigh = 1; %ace high encoding
10 numGames = 200000; %number of simulation games
   %Initialize win counters
12 trainedPolicyWin = zeros(1,numGames);
   randomPolicyWin = zeros(1,numGames);
14 fixedPolicyWin = zeros(1,numGames);
   for i = 1:numGames
16     handOver = 0;
       win = [];
18     %Deal cards to numCards x numPlayers matrix
       [dealtCards,deck] = dealCards(numPlayers,numCards,numDecks
           [+],aceHigh);
20     %Split cards into player and dealer hands
       pHand = dealtCards(1,:); dHand = dealtCards(2,:);
22     shownCard = dHand(2);
       %Evaluate player hand (triplicate, one for each player)
24     [pHand,pVal,pBust,pUseAce] = HandEval(pHand);
       %Store evaluations in array to evaluate each player
           [+],separately by idx
26     Hand = [pHand; pHand; pHand];
       Val = [pVal; pVal; pVal];
28     Bust = [pBust; pBust; pBust];
       Ace = [pUseAce; pUseAce; pUseAce];
30     Ace = Ace + 1;
       deck_reuse = deck;
32     for plyr = 1:3
           hand = Hand(plyr,:);
34         handOver = 0;
           deck = deck_reuse;
36         while ~Bust(plyr) && ~handOver
             act = ChooseAction(hand,Val(plyr),shownCard,Ace(
                 [+],plyr),plyr,Q);
38             if act == 2
                 [hand,deck] = PlayerHit(hand,deck);
40             [hand,Val(plyr),Bust(plyr),Ace(plyr)] =

```

```

        [+]HandEval(hand);
    else
        handOver = 1;
    end
end
end
[dBust,dBJ,dVal,deck] = DealerPolicy(dHand,deck);
for plyr = 1:3
    %Determine winner
    if (Bust(plyr)) || ((dVal > Val(plyr)) && (~dBust))
        win = [win 0];

    elseif (Val(plyr) > dVal) || dBust
        win = [win 1];
    else
        win = [win 0];
    end
end
trainedPolicyWin(i) = win(1);
randomPolicyWin(i) = win(2);
fixedPolicyWin(i) = win(3);
end
%Calculate win % for each agent
tWinRate = sum(trainedPolicyWin)/length(trainedPolicyWin)
rWinRate = sum(randomPolicyWin)/length(randomPolicyWin)
fWinRate = sum(fixedPolicyWin)/length(fixedPolicyWin)

```

- Simulation Select Action Function

```

function [ action ] = ChooseAction( hand,handVal,shownCard,
    [+]uAce,agent_type,Q_table )
2 %function chooses agent actions (hit,stay) for AgentAnalytics.
    [+]m
    %Trained agent
    if agent_type == 1
        action = max(Q_table(handVal,shownCard,uAce,:));
    %Random agent
    elseif agent_type == 2
        actRand = rand;
        if actRand < 0.5
            action = 1;
        else
            action = 2; %action = 2;
        end
    end
end

```

```

14         %Fixed policy agent
15     else
16         if handVal < 17
17             action = 2;
18         else
19             action = 1; %action = 2;
20         end
21     end
22 end

```

5.2 Deep Q-Learning Listing

```

1  #DEEP Q-LEARNING FOR BLACKJACK
   #AUTHOR: Ian McAtee
3  #DATE: 12/06/2021
   #CLASS: EE5885 Deep-Reinforcement Learning
5  #DESCRIPTION: Script to perform the deep Q-learning for an agent
   [+]to play blackjack

7  #DQN AGENT CLASS

9  import random
   import numpy as np
11 from keras.models import Sequential
   from keras.layers.core import Dense, Activation
13 from tensorflow.keras.optimizers import Adam

15 class DQNAgent():
    def __init__(self, env, epsilon=1.0, alpha=0.5, gamma=0.9):
17
        self.action_size = 2 #Size of action space
19        self.state_size = env.observation_space #Size of state
            [+]space
        self.epsilon = epsilon #Exploartion
            [+]factor
21        self.alpha = alpha #Learning Rate
        self.gamma = gamma #Discount Factor
23        self.model = self.QNetwork() #QNetwork
            [+]Approximation
        self.target = self.QNetwork() #Target Network

25        #Construct Q-function approximation network
27        def QNetwork(self):
            model = Sequential()

```

```

29     model.add(Dense(32, input_shape = (2,), kernel_initializer
        [=]'random_uniform', activation='relu'))
        model.add(Dense(16, activation='relu'))
31     model.add(Dense(self.action_size, activation='softmax'))
        model.compile(loss='mse', optimizer=Adam(learning_rate=
            [=]self.alpha))
33
        return model
35
#Function to choose action based on epsilon greedy policy
37 def chooseAction(self, state):

39     #Random Action
        if random.random() < self.epsilon:
41         action = random.choice([0,1])

43     #Greedy Action
        else:
45         action_value = self.model.predict(state)
            action = np.argmax(action_value[0])
47     return action

49 #Function to do the deep Q-learning
def DQlearning(self, state, action, reward, next_state, done,
[+]steps):

51     #Set the target

53
        #If done with hand set target to the reward (because there
            [=] is no next state)
55     if done:
        y = reward

57
        #If not done with hand set target as in Q-learning
            [=]equation, uses target network
59     else:
        y = reward + self.gamma * np.amax(self.target.predict(
            [=]next_state)[0])

61
        #Use Q-network approximation to predict the Q-function for
            [=] current state
63     Q_sa = self.model.predict(state)

65
        #Set the current action's index in the predicted Q-
            [=]function to the target

```

```

        Q_sa[0][action] = y
67
        #Train model
69        self.model.fit(state, Q_sa, epochs=1, verbose=0)

71        #M is number of steps before updating target network
        M = 200;

73
        if steps%M == 0:
75            self.updateTarget() #Update target network

77        #Function to update the target network
        def updateTarget(self):
79            self.target.set_weights(self.model.get_weights())

81        #Function to evaluate trained QNetwork given a state sample
        def evaluateModel(self, state):
83            pred = max(self.model.predict(state))
            maxQ = max(pred)
85            if pred[0] == maxQ:
                a = 0
87            else:
                a = 1
89            return a

91
#DEEP Q-LEARNING TRAINING
93
import gym
95 import matplotlib.pyplot as plt

97 env = gym.make('Blackjack-v0')
    env.seed(0)
99
    #Number of hands to play in each episode
101 numSamps = 50

103 #Number of episodes to play
    numEpisodes = 50
105
    #Create deep Q learning Agent
107 agent = DQNAgent(env=env, epsilon=0.1, alpha=0.001, gamma=0.9)

109 #Define the average reward
    averageReward = []

```

```

111 #Initialize a step counter
113 step = 1

115 #Setup the environment and reshape the state space for neural
    [+]network
    state = env.reset()
117 state = np.reshape(state[0:2], [1,2])

119 #Perform the Deep Q-learning
    for episode in range(numEpisodes): #Loop Episodes
121
        print('Training: Episode ',episode)
123
        totalReward = 0
125
        for samp in range(numSamps): #Loop samples
127
            #Take action based on epsilon-greedy policy
129            action = agent.chooseAction(state)
            #Apply action to environment to get next state, reward,
            [+]and is_done signal
131            next_state, reward, done, _ = env.step(action)
            #Reshape state space to use with neural network
133            next_state = np.reshape(next_state[0:2], [1,2])
            #Tally the reward
135            totalReward += reward
            #Perform the deep Q-learning
137            agent.DQlearning(state, action, reward, next_state, done,
                [+]step)
            #Go to next state
139            state = next_state
            state = np.reshape(state[0:2], [1,2])
141            #If blackjack hand done, reset the environment
            if done:
143                state = env.reset() # Environment deals new cards to
                    [+]player and dealer
                    state = np.reshape(state[0:2], [1,2])
145
            step += 1
147
        averageReward.append(totalReward)
149
151

```

```

#Plot the average rewards over the episodes
153 plt.plot(averageReward)
    plt.xlabel('Episodes')
155 plt.ylabel('Average Reward')
    plt.show()

157
#DETERMINE POLICY
159
#Get policy table, by evaluating Q-funtion approximation network
161 possHandSum = range(4, 22)
    possDealUpCard = range(1, 11)
163 policyTable = [[str(-1) for i in possDealUpCard] for j in
    [+]possHandSum]
    for i in possHandSum:
165         for j in possDealUpCard:
            policyTable[i-4][j-1] = str(agent.evaluateModel([[i,j]]))

167
#Display the Policy
169 print('
                                POLICY ')
    print(' ')
171 print('Agent Hand |
                                Dealer Up Card')
    print(' Sum | A 2 3 4 5 6 7 8 9
    [+] 10 ')
173 print('
    [+]-----
    [+]')
    for i in range(4,10):
175         print(' ',i, ' | ',policyTable[i-4])
    for i in range(10,22):
177         print(' ',i, ' | ',policyTable[i-4])

179 #TEST THE FOUND POLICY AGAINST THE HEURISTIC POLICY

181 #Resetl the environment
    env.seed(0)
183 state = env.reset()

185
#Test Heuristic Hit Under 17 Policy
187 numHands = 1000
    winsH = 0

189
    for i in range(numHands):
191         state = env.reset()
            done = False

```

```

193     while not done:
194
195         #Hit or stay based on if you hold less than 17
196         if state[0] < 17:
197             action = 1
198         else:
199             action = 0
200
201         #Get the next state
202         next_state, reward, done, _ = env.step(action)
203
204         state = next_state
205
206         if reward == 1.0:
207             winsH += 1
208
209 #Test Trained Q-Learning Policy
210 numHands = 1000
211 winsQ = 0
212
213 for i in range(numHands):
214     state = env.reset()
215     done = False
216     while not done:
217
218         #Evaluate the q-function approximation network
219         action = agent.evaluateModel([[state[0], state[1]]])
220
221         #Get the next state
222         next_state, reward, done, _ = env.step(action)
223
224         state = next_state
225
226         if reward == 1.0:
227             winsQ += 1
228
229 #Display the win results
230 print('Heuristic Wins: ', winsH)
231 print('Qlearning Wins: ', winsQ)

```

References

- [1] E. Thorpe, "A Favorable Strategy for Twenty-One. *Proc Natl Acad Sci USA*, vol. 47, 1, 1961